



भारतीय प्रौद्योगिकी संस्थान हैदराबाद
Indian Institute of Technology Hyderabad

EE6310, Image and Video Processing

Image Analysis Techniques

February 14, 2023

Image Analysis Techniques

- ▶ Image Quality Assessment
- ▶ Edge Detection
- ▶ Hough Transform
- ▶ Template matching, SIFT

Image Quality Assessment

- ▶ Big question: what determines image **quality**?
- ▶ Ultimate receiver is the **human eye – subjective judgment** is all that matters
- ▶ How to assess quality **objectively** – i.e., by an **algorithm**?

Image Quality Assessment – Subjective Quality



Figure: One method!

Image Quality Assessment – Classification

Techniques can be broadly classified into:

- ▶ Full-reference: **both** test image and reference image available (discussed in this class)
- ▶ Reduced reference: test image and **certain characteristics** of reference/distortion available
- ▶ No reference: **only** test image available

An **active** research area with **immediate practical** applications!

Image Quality Assessment – MSE/PSNR

- ▶ The **mean squared error** is the long-standing most widely used IQA method
- ▶ Definition: given **reference** image I and **test/observed** image J of size $M \times N$,

$$MSE(I, J) = \frac{1}{NM} \sum_{i=0}^{N-1} \sum_{j=0}^{M-1} [I(i, j) - J(i, j)]^2$$

- ▶ **Peak Signal to Noise Ratio (PSNR)** is defined as:
 $PSNR(I, J) = 10\log_{10} \frac{L^2}{MSE(I, J)}$ where L is the range of allowable gray values

Image Quality Assessment – MSE/PSNR Pros and Cons

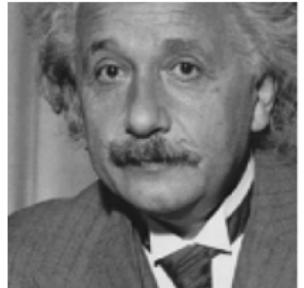
Advantages

- ▶ **Computationally simple**
- ▶ Analytically easy to work with
- ▶ Easy to **optimize** w.r.t. MSE
- ▶ **Effective** in high SNR situations

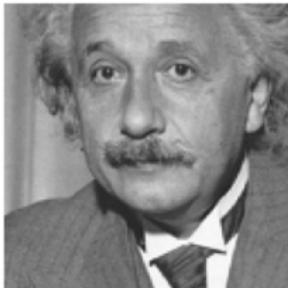
Disadvantages

- ▶ **Poor** correlation with human visual perception!

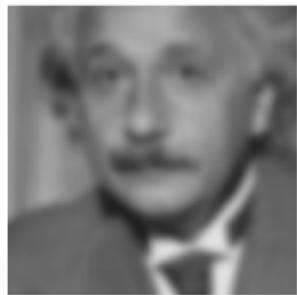
Image Quality Assessment – MSE/PSNR Failure



(a) Reference



(b) Mean



(c) Blurred



(d) JPEG

Figure: All images have same MSE (309)!

Image Quality Assessment – Human Subjectivity

- ▶ **Human opinion** is the ultimate gauge of image quality
- ▶ **Calibrated test conditions** required to **measure true image quality**
- ▶ The resulting **mean opinion score (MOS)** correlated with QA algorithm performance
- ▶ Until recently, MSE had **little competition** despite 40 years of research!

Image Quality Assessment – Structural Similarity Index

- ▶ Modern, successful approach: measure **loss of structure** in a distorted image
- ▶ Basic idea: combine local measures of similarity of **luminance, contrast, structure** into local measure of quality
- ▶ Perform **weighted average** of local measure across image

Image Quality Assessment – Structural Similarity Index

The structural similarity (SSIM) index between I and J at a point (i,j) is expressed as

$$\text{SSIM}_{I,J}(i,j) = L_{I,J}(i,j)C_{I,J}(i,j)S_{I,J}(i,j) \text{ where}$$

- ▶ $L_{I,J}(i,j)$ is a measure of local **luminance similarity**
- ▶ $C_{I,J}(i,j)$ is a measure of local **contrast similarity**
- ▶ $S_{I,J}(i,j)$ is a measure of local **structure similarity**

Image Quality Assessment – Luminance Similarity

Luminance similarity is defined as

$$L_{I,J}(i,j) = \frac{2\mu_I(i,j)\mu_J(i,j)+C_1}{\mu_I(i,j)^2+\mu_J(i,j)^2+C_1} = \frac{2\mu_I\mu_J+C_1}{\mu_I^2+\mu_J^2+C_1}$$

where

$$\mu_I(i,j) = \sum_{p=-P}^P \sum_{q=-Q}^Q w(p,q)I(i+p, j+q),$$

$$\sum_{p=-P}^P \sum_{q=-Q}^Q w(p,q) = 1.$$

$w(p, q)$ is an **isotropic, unit area weighting function** and C_1 is a stabilizing constant

Image Quality Assessment – Contrast Similarity

Contrast similarity is defined as

$$C_{I,J}(i,j) = \frac{2\sigma_I(i,j)\sigma_J(i,j)+C_2}{\sigma_I(i,j)^2+\sigma_J(i,j)^2+C_2} = \frac{2\sigma_I\sigma_J+C_2}{\sigma_I^2+\sigma_J^2+C_2}$$

where

$$\sigma_I(i,j) = \sqrt{\sum_{p=-P}^P \sum_{q=-Q}^Q w(p,q)[I(i+p,j+q) - \mu_I(i,j)]^2},$$

C_2 is a stabilizing constant

Image Quality Assessment – Structural Similarity

Structural similarity is defined as

$$S_{I,J}(i,j) = \frac{\sigma_{I,J}(i,j) + C_3}{\sigma_I(i,j)\sigma_J(i,j) + C_3} = \frac{\sigma_{I,J} + C_3}{\sigma_I\sigma_J + C_3}$$

where

$$\sigma_{I,J}(i,j) =$$

$$\sum_{p=-P}^P \sum_{q=-Q}^Q w(p,q)[I(i+p, j+q) - \mu_I(i, j)][J(i+p, j+q) - \mu_J(i, j)],$$

C_3 is a stabilizing constant

Image Quality Assessment – SSIM Properties

- ▶ Symmetry: $\text{SSIM}_{I,J}(i,j) = \text{SSIM}_{J,I}(i,j)$
- ▶ Boundedness: $-1 \leq \text{SSIM}(I, J) \leq 1$
- ▶ Unique maximum: $\text{SSIM}(i,j) = 1$ iff I, J are locally identical:
 $I \diamond w = J \diamond w$
- ▶ C_1, C_2, C_3 are **stabilizers** in case the local means or contrasts
are **very small**
- ▶ For 8-bit gray scale images,
 $C_1 = (0.01 * 255)^2, C_2 = (0.03 * 255)^2, C_3 = C_2/2$ work well

Image Quality Assessment – SSIM Map

- ▶ Displaying $\text{SSIM}(i,j)$ as an image is called an **SSIM Map**. It is an effective way of **visualizing** where the images I, J differ
- ▶ The **SSIM map** depicts where the quality of one image differs from the other

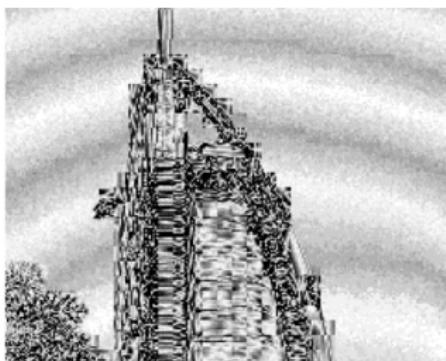
Image Quality Assessment – SSIM Map Example



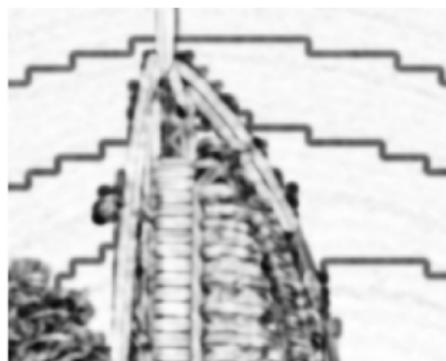
(a)



(b)



(c)



(d)

Figure: a: Reference; b: JPEG; c: Absolute diff; d: SSIM map

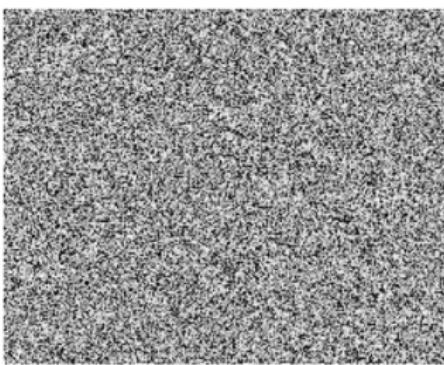
Image Quality Assessment – SSIM Map Example



(a)



(b)



(c)



(d)

Figure: a: Reference; b: AWGN; c: Absolute diff; d: SSIM map

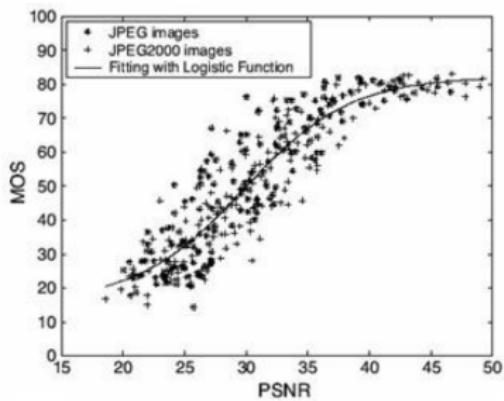
Image Quality Assessment – SSIM Map

- ▶ The mean SSIM is the average value of $\text{SSIM}(i,j)$ over the entire image:

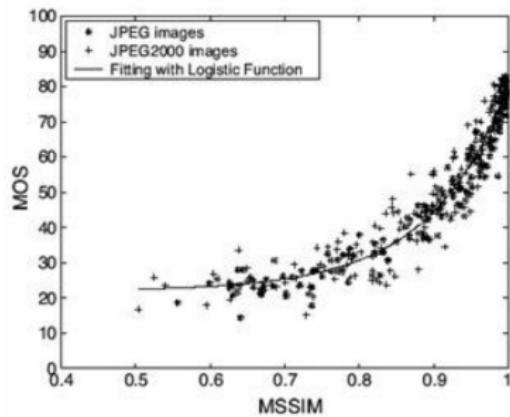
$$\text{SSIM}(I, J) = \frac{1}{NM} \sum_{i=0}^{N-1} \sum_{j=0}^{M-1} \text{SSIM}_{I,J}(i,j)$$

- ▶ The mean SSIM correlates extraordinarily well with human response as measured by mean opinion score (MOS)

Image Quality Assessment – SSIM vs PSNR

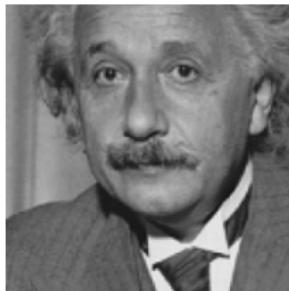


(a) PSNR vs MOS

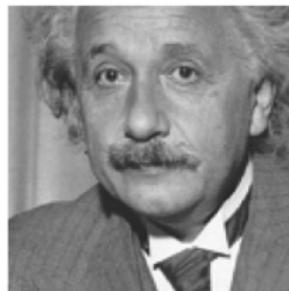


(b) MSSIM vs MOS

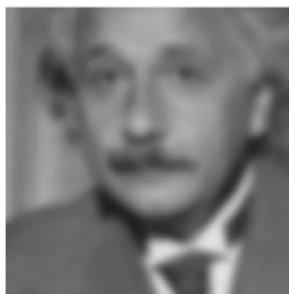
Image Quality Assessment – MSSIM Example



(a) Reference,
MSSIM = 1



(b) Mean, MSSIM
= 0.91



(c) Blurred,
MSSIM = 0.64

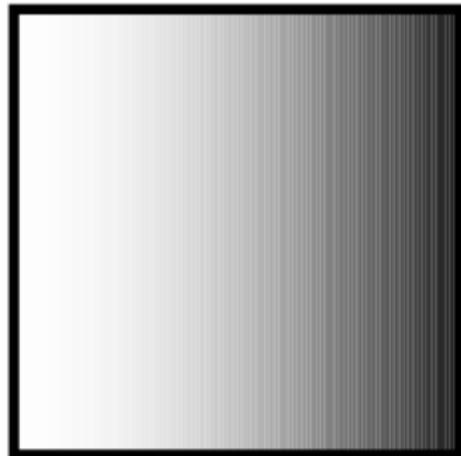


(d) JPEG, MSSIM
= 0.58

Figure: All images have same MSE (309)!

Edge Detection – Edge Examples

Where are the edges?



Edge Detection – Introduction

- ▶ Probably the **most exhaustively** studied topic in image analysis
- ▶ **Edges** are sudden, sustained changes in **AVERAGE** image intensity that extent along a contour
- ▶ **Edge detection** used as a precursor to most practical image analysis tasks
- ▶ Reasons:
 - ▶ **Enormous information** contained in image edges. Image can largely be recognized from edges alone!
 - ▶ An **edge map** requires **much less** storage than the image itself – it is a **binary** contour plot

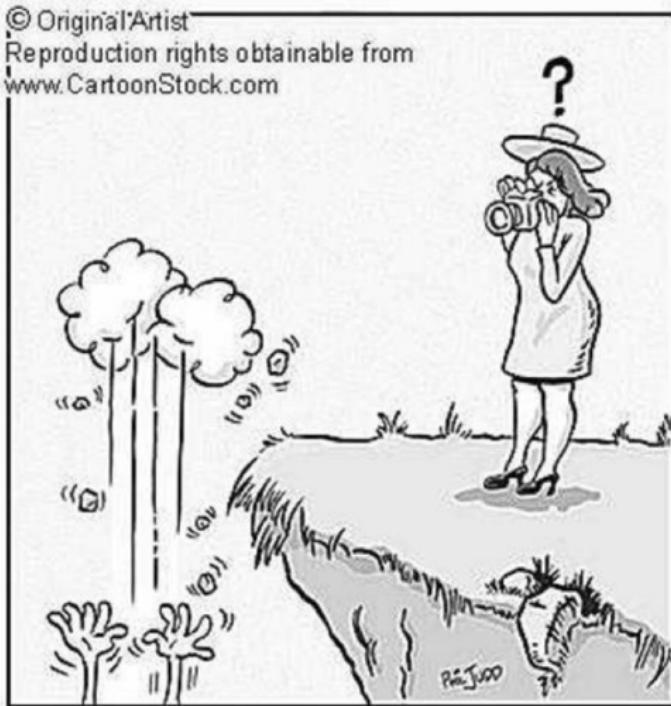
Edge Detection – Overview of Methods

- ▶ Methods for edge detection studied since mid-1960s
- ▶ Easily the most studied problem in image analysis!
- ▶ **Hundreds** of approaches exist – based on a myriad of **math techniques**
- ▶ Problem became well understood in the 1980s
- ▶ In this class, we'll study **three fundamental methods:**
 - ▶ **Gradient** edge detectors
 - ▶ **Laplacian** edge detectors
 - ▶ **Diffusion-based** edge detectors

Edge Detection – Before Computers!

© Original Artist

Reproduction rights obtainable from
www.CartoonStock.com



"Move back just a little Fred....Fred!?!"

Edge Detection – Gradient Edge Detectors

- ▶ Oldest (Roberts 1965) but still very valuable class of edge detectors
- ▶ For a 2-D continuous function $f(x, y)$, the **gradient** is a two-element **vector**: $\nabla f(x, y) = [f_x(x, y), f_y(x, y)]^T$ where
 - ▶ $f_x(x, y) = \frac{\partial}{\partial x} f(x, y)$,
 - ▶ $f_y(x, y) = \frac{\partial}{\partial y} f(x, y)$are the **directional derivatives** of $f(x, y)$ along the x and y directions respectively

Edge Detection – Gradient Measurements

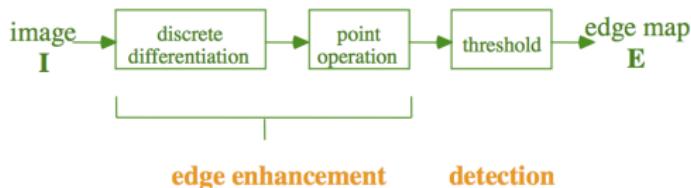
- ▶ The **direction** of the **fastest rate of change** of $f(x, y)$ at a point (x, y) is the **gradient orientation**
$$\theta_f(x, y) = \tan^{-1}\left(\frac{f_y(x, y)}{f_x(x, y)}\right)$$
- ▶ The rate of change is the **gradient magnitude**
$$M_f(x, y) = \sqrt{f_x(x, y)^2 + f_y(x, y)^2}$$
- ▶ Gradient is appealing since we can expect edges to locally exhibit the **greatest rate of change** of what?

Edge Detection – Isotropic Gradient Magnitude

- ▶ Taking derivative in **any two perpendicular directions**, say (x', y') **does not change** $M_f(x, y)$
- ▶ $M_f(x, y)$ is **rotationally symmetric** or **isotropic**
- ▶ **Isotropicity** desirable since we want to detect edges regardless of orientation

Edge Detection – Gradient Edge Detector Diagram

Described by the following block diagram:



- ▶ **Digital differentiation:** digitally approximate ∇I
- ▶ **Point operation:** estimate $M_I = |\nabla I|$
- ▶ **Threshold:** decide which large values of M_I are likely locations

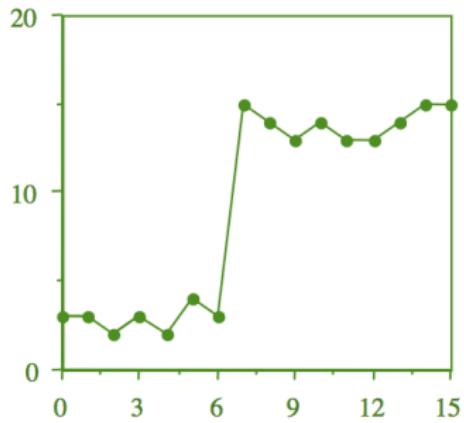
Edge Detection – Digital Differentiation

- ▶ Digital differentiation is **differencing**. For a 1-D function $f(x)$ that has been sampled to produce $f(i)$, either:
 - ▶ $\frac{d}{dx}f(x)|_{x=i} \approx f(i) - f(i-1)$ or
 - ▶ $\frac{d}{dx}f(x)|_{x=i} \approx \frac{f(i+1) - f(i-1)}{2}$
- ▶ Advantage of first method is that the current value $f(i)$ is used in computation
- ▶ Advantage of second method is that it is **centered**

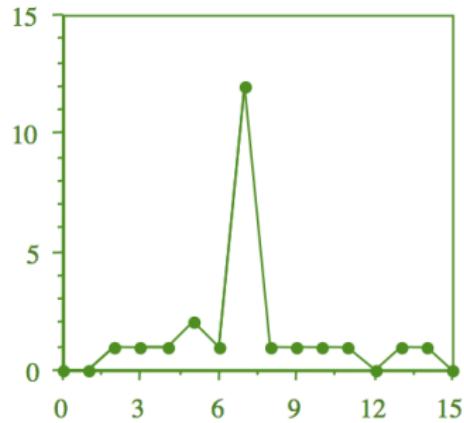
Edge Detection – 2-D Differentiating

- ▶ Similarly extended to **2-D**
- ▶ Current
 - ▶ $\frac{\partial}{\partial x} f(x, y)|_{(x,y)=(i,j)} \approx f(i, j) - f(i - 1, j)$ or
 - ▶ $\frac{\partial}{\partial y} f(x, y)|_{(x,y)=(i,j)} \approx f(i, j) - f(i, j - 1)$
- ▶ Centered
 - ▶ $\frac{\partial}{\partial x} f(x, y)|_{(x,y)=(i,j)} \approx \frac{f(i+1,j) - f(i-1,j)}{2}$ or
 - ▶ $\frac{\partial}{\partial y} f(x, y)|_{(x,y)=(i,j)} \approx \frac{f(i,j+1) - f(i,j-1)}{2}$

Edge Detection – Example 1-D Differencing

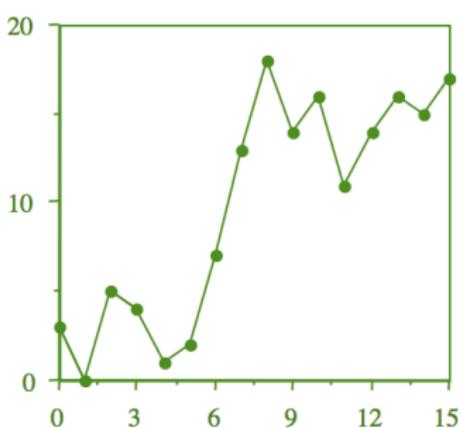


(a) $I(i)$

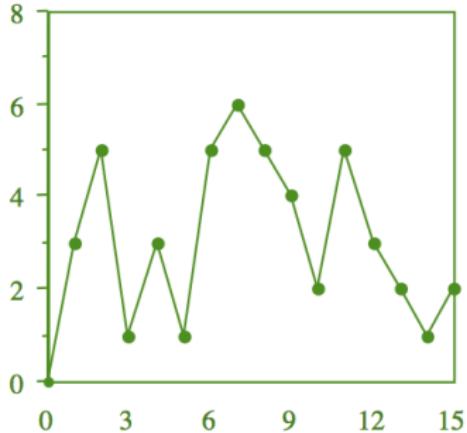


(b) $D_I(i) = |I(i) - I(i - 1)|$

Edge Detection – Example 1-D Differencing with Noise



$$(a) \quad J(i) = I(i) + N(i)$$



$$(b) \quad D_J(i) = |J(i) - J(i - 1)|$$

Noise is a **huge** problem because differentiation **always** emphasises high frequencies (like noise)

Edge Detection – Gradient Edge Detectors

Define **convolutional edge templates** Δ_x and Δ_y that produce directional derivative estimates:

- ▶ Adjacent: $\Delta_x = \begin{bmatrix} -1 & +1 \end{bmatrix}$, $\Delta_y = \begin{bmatrix} -1 \\ +1 \end{bmatrix}$
- ▶ Centered: $\Delta_x = \begin{bmatrix} -1 & 0 & +1 \end{bmatrix}/2$, $\Delta_y = \begin{bmatrix} -1 \\ 0 \\ +1 \end{bmatrix}/2$
- ▶ Roberts: $\Delta_x = \begin{bmatrix} -1 & 0 \\ 0 & +1 \end{bmatrix}$, $\Delta_y = \begin{bmatrix} 0 & -1 \\ +1 & 0 \end{bmatrix}$

All three give very **similar** performance.

Edge Detection – Noise Reducing Detectors

Designed to reduce noise effects by averaging along rows and columns:

► Prewitt: $\Delta_x = \begin{bmatrix} -1 & 0 & +1 \\ -1 & 0 & +1 \\ -1 & 0 & +1 \end{bmatrix} / 3$, $\Delta_y = \begin{bmatrix} -1 & -1 & -1 \\ 0 & 0 & 0 \\ +1 & +1 & +1 \end{bmatrix} / 3$

► Sobel: $\Delta_x = \begin{bmatrix} -1 & 0 & +1 \\ -2 & 0 & +2 \\ -1 & 0 & +1 \end{bmatrix} / 4$, $\Delta_y = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ +1 & +2 & +1 \end{bmatrix} / 4$

Edge Detection – Gradient Magnitude

- ▶ The **point operation** combines directional derivative estimates Δ_x and Δ_y into a single estimate of **gradient estimate**
- ▶ The usual estimates:
 - ▶ A: $M(i,j) = \sqrt{\Delta_x^2(i,j) + \Delta_y^2(i,j)}$
 - ▶ B: $M(i,j) = |\Delta_x(i,j)| + |\Delta_y(i,j)|$
 - ▶ C: $M(i,j) = \max\{|\Delta_x(i,j)|, |\Delta_y(i,j)|\}$
- ▶ The following always holds: $C \leq A \leq B$
- ▶ A is the **correct** interpretation but B, C are **cheaper** – no squaring/root operations
- ▶ B often **overestimates** and C often **underestimates** edge magnitude

Edge Detection – Gradient Magnitude

- ▶ D: $M(i,j) = \max\{|\Delta_x(i,j)|, |\Delta_y(i,j)|\} + \frac{1}{4}\min\{|\Delta_x(i,j)|, |\Delta_y(i,j)|\}$
- ▶ D better than B or C. **Slight** difference between A and D

Edge Detection – Magnitude Thresholding

- ▶ Once edge magnitude is computed, it must be **thresholded** to find plausible edge locations
- ▶ This produces binary **edge map E**:
$$E(i,j) = \begin{cases} 1; & M(i,j) > \tau \\ 0; & M(i,j) \leq \tau \end{cases}$$
- ▶ Thus:
 - ▶ '1': edge **present** at (i,j)
 - ▶ '0': edge **absent** at (i,j)
- ▶ Threshold τ constrains the **sharpness** and **magnitude** of edges that are detected

Edge Detection – Gradient Detection Advantages

- ▶ Simple, computationally efficient
- ▶ Natural definition
- ▶ Works well on “clean” images

Edge Detection – Gradient Detection Disadvantages

- ▶ Extremely **noise-sensitive**
- ▶ Requires a **threshold** – difficult to select – usually requires **interactive selection**
- ▶ Gradient magnitude estimate **falls above threshold** over a few pixels distance from the true edge. So edges are **a few pixels wide**
- ▶ This usually requires some sort of “**edge thinning**” operation - usually heuristic
- ▶ The edge contours are often **broken** – gaps appear. This requires an “**edge linking**” operation – usually heuristic

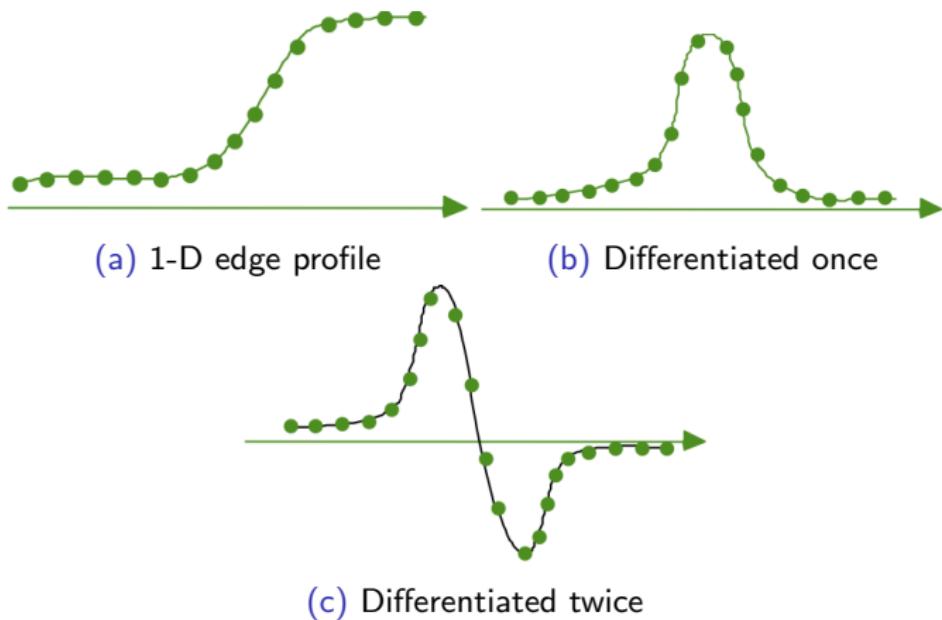
Edge Detection – Laplacian Edge Detectors

- ▶ Edge detectors are based on **second derivates**
- ▶ For a continuous 2-D function $F(x, y)$, the **Laplacian** is defined as: $\nabla^2 f(x, y) = \frac{\partial^2}{\partial x^2} f(x, y) + \frac{\partial^2}{\partial y^2} f(x, y)$
- ▶ It is a **scalar**
- ▶ **Fact:** Laplacian remains **unchanged** if directional derivates taken w.r.t. **any other set of orthogonal directions**

Edge Detection – Laplacian Edge Detectors

- ▶ The **Laplacian edge detection** process can be broken down into:
 - ▶ Discrete differentiation that causes **edge enhancement**
 - ▶ Zero-crossing detection to **detect** edges
- ▶ **Digital differentiation:** Digitally approximate $\nabla^2 I$
- ▶ **Zero-crossing detection:** Discover where the Laplacian crosses the zero level

Edge Detection – Laplacian Edge Detectors



A **zero crossing** or **ZC** occurs near the center of the edge where the **slope of the slope changes sign**

Edge Detection – Digital Twice Difference

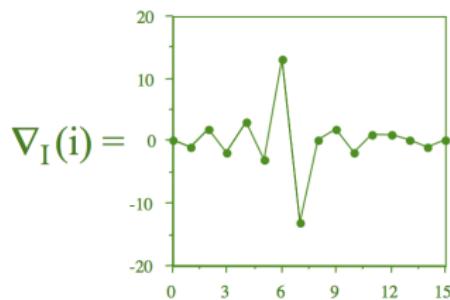
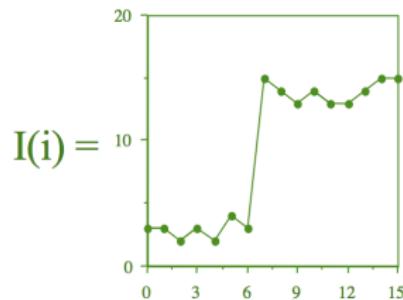
- ▶ For a 1-D function $f(x)$ we use:
 $\frac{d}{dx} f(x)|_{x=i} \approx f(i) - f(i-1) = y(i)$
- ▶ $\frac{d^2}{dx^2} f(x)|_{x=1} \approx y(i) - y(i-1) = f(i+1) - 2f(i) + f(i-1)$
with the **convolution template**:
$$\begin{bmatrix} 1 & -2 & 1 \end{bmatrix}$$
- ▶ In 2 dimensions: $\nabla^2 I(x,y) \approx [I(i+1,j) - 2I(i,j) + I(i-1,j)] + [I(i,j+1) - 2I(i,j) + I(i,j-1)]$

- ▶ **Convolution template:**
$$\begin{bmatrix} 1 & -2 & 1 \end{bmatrix} + \begin{bmatrix} 1 \\ -2 \\ 1 \end{bmatrix} =$$

$$\begin{bmatrix} 0 & +1 & 0 \\ +1 & -4 & +1 \\ 0 & +1 & 0 \end{bmatrix}$$

Edge Detection – Digital Twice Difference Example

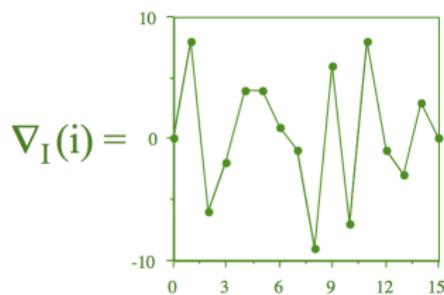
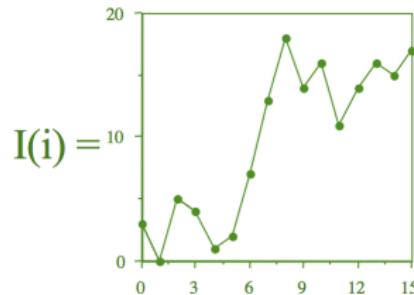
► $\nabla_I(i) = I(i+1) - 2I(i) + I(i-1)$



- Clearly reveals a sharp edge location: a single **large slope ZC** and several smaller ZCs

Edge Detection – Digital Twice Difference Example

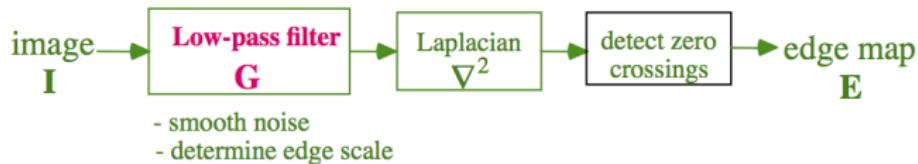
► $I(i,j) = J(i,j) + N(i,j)$



- Numerous **spurious ZCs**
- **Noise an even bigger problem.** Differentiating twice creates **highly amplified noise**

Edge Detection – Smoothed/Multi-scale Laplacian Edge Detection

- ▶ Laplacian operator too noise sensitive to be practical
- ▶ A simple modification makes it very powerful
- ▶ Basic idea:



- ▶ Difference: a **linear blur** applied prior to Laplacian operator

Edge Detection – Low pass pre-filter

- ▶ The **main purpose** of a low pass pre-filter to the Laplacian is to **attenuate high-frequency noise** while retaining significant image structure
- ▶ The **secondary purpose** of smoothing is to **constrain the scale** over which edges are detected
- ▶ A **high pass** operation followed by a **low pass** operation results in a **band pass** filter if their passbands overlap

Edge Detection – Gaussian pre-filter

- ▶ A lot of research has been done on how to select the low pass filter G
- ▶ It has been found that the optimal smoothing filter in the following two senses:
 1. best **edge location accuracy**
 2. **maximum** signal-to-noise ratio (SNR)

is a **Gaussian** filter: (K is an irrelevant constant)

$$G(i,j) = K \cdot \exp\{-[i^2 + j^2]/2\sigma^2\}$$

Edge Detection – Laplacian of Gaussian Edge Detector

- ▶ Define the **Laplacian of a Gaussian** or **LoG** on I :
$$J(i,j) = \nabla^2[G(i,j) * I(i,j)] = G(i,j) * \nabla^2I(i,j) =$$
$$\nabla^2G(i,j) * I(i,j)$$
- ▶ Above 3 forms equivalent since **linear operations commute**
- ▶ Best approach: pre-compute **LoG**:
$$\nabla^2G(i,j) \propto [1 - \frac{i^2+j^2}{\sigma^2}].\exp[-\frac{i^2+j^2}{2\sigma^2}]$$
 and convolve with image I
- ▶ Constant multiplier omitted since only ZC of interest to us

Edge Detection – Polar Form of LoG

- The LoG is **isotropic** and can be written in polar form as:
$$\nabla^2 G(r) \propto (1 - \frac{r^2}{\sigma^2}) \cdot \exp(-\frac{r^2}{2\sigma^2})$$

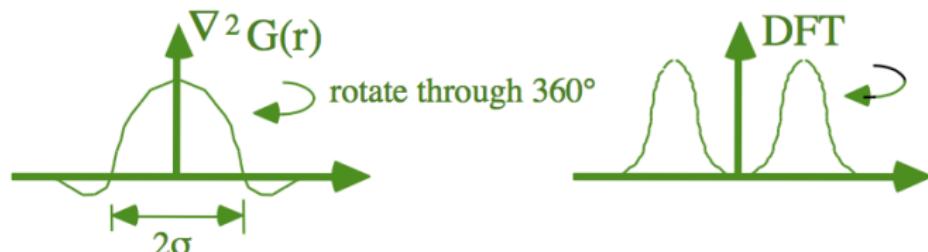
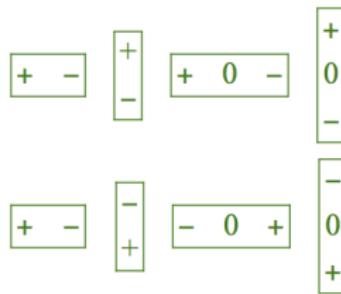


Figure: LoG in space and frequency

Edge Detection – ZC Detection

- ▶ The last stage of edge detection is **zero-crossing detection**
- ▶ Let $J = J(i, j)$ be the result of LoG filtering
- ▶ A ZC is a **crossing** of the zero level: the algorithm must search for pixel occurrences of the form:



- ▶ By convention, one sign or the other is **marked** as the edge unless finer resolution (sub-pixel) is required

Edge Detection – Scale of LoG

- ▶ The larger the value of σ used, the greater the degree of **smoothing** by the low-pass pre-filter G
- ▶ If σ is **large**, then noise will be greatly smoothed - but so will **less significant edges**
- ▶ **Noise sensitivity increases** with decrease in σ but LoG detector then detects more detail

Edge Detection – Digital Implementation of LoG

- ▶ Use the **sampled LoG**:

$$\nabla^2 G(i,j) \propto [1 - \frac{i^2+j^2}{\sigma^2}] \cdot \exp[-\frac{i^2+j^2}{2\sigma^2}]$$

- ▶ Specific rules of thumb should be followed
- ▶ **Enough** of $\nabla^2 G(i,j)$ must be sampled. LoG will not work unless the template contains both **main** and **minor lobes**
- ▶ In practice, the **radius** R of the LoG in space should satisfy $R \geq 4\sigma$ (in pixels)
- ▶ Adjust template to **sum to zero** by subtracting mean – why?
- ▶ LoG will not work unless $\sigma \geq 1$

Edge Detection – Advantages of LoG

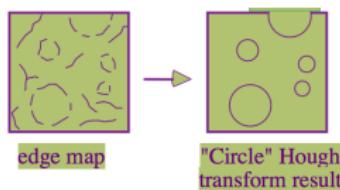
- ▶ Usually **doesn't require thresholding**
- ▶ Yields **single pixel-width edges always**
- ▶ Yields **connected edges always**
- ▶ Appears close to what happens in **biological vision**

Edge Detection – Disadvantages of LoG

- ▶ **More computation** than gradient edge detectors
- ▶ ZCs continuity property often leads to ZCs that **meander across the image**
- ▶ ZC contours tend to be **over-smooth near corners**

Hough Transform

- ▶ The **Hough Transform** is a simple, generalizable tool for finding **instances of curves of a specific shape** in a **binary edge map**
- ▶ Pictorially:



Hough Transform – Advantages

- ▶ Highly **noise-insensitive** - can “pick” edges from among many spurious edges
- ▶ Able to **reconstruct “partial” curves** containing gaps and breaks to the “ideal” form
- ▶ Can be **generalized** to **almost any desired shape**

Hough Transform – Disadvantages

- ▶ It is **computation** and **memory-intensive**

Basic Hough Transform

- ▶ Assume that it is desired to find the locations of curves that
 - ▶ Can be expressed as a function of (i, j)
 - ▶ Have a set (vector) of parameters $\mathbf{a} = [a_1, \dots, a_n]^T$ that specify the exact **size**, **shape** and **location** of the curves
- ▶ Thus curves of the form: $f(i, j; \mathbf{a}) = 0$

Basic Hough Transform – Curves With Parameters

- ▶ 2-D **lines** have a **slope-intercept** form

$f(i, j; \mathbf{a}) = j - mi - b = 0$, where $\mathbf{a} = (m, b) = (\text{slope}, j\text{-intercept})$

- ▶ 2-D **circles** have the form

$f(i, j; \mathbf{a}) = (i - i_0)^2 + (j - j_0)^2 - r^2 = 0$, where $\mathbf{a} = (i_0, j_0, r) = (\text{center coordinates, radius})$

Basic Hough Transform

- ▶ **Input** to Hough Transform: An **edge map** or other representation of **image contour**
- ▶ From this map, a **record** is made of **possible instances** of the shape in the image
- ▶ The likelihood of a shape is found by **counting the number of edge points that fall on the shape countour**

Hough Example

- ▶ Straight line/line segment detection



- ▶ The line indicated is the “most likely” since there are **seven** pixels that contribute **evidence** for its existence
- ▶ There are other “less likely” segments

Hough Accumulator

- ▶ The **Hough accumulator** \mathbf{A} is a **matrix** that **accumulates evidence** of instances of the curve of interest via **counting**
- ▶ The Hough accumulator \mathbf{A} is **n-dimension**, where n is the number of **parameters** in \mathbf{a} : $\mathbf{a} = [a_1, a_2, \dots, a_n]^T$
- ▶ Each parameter $a_i, i = 1, \dots, n$ can take on only a **finite** number of values N_i
- ▶ The accumulator \mathbf{A} is an $N_1 \times N_2 \times \dots N_{n-1} \times N_n$ matrix containing $N_1.N_2.\dots.N_{n-1}.N_n$ slots

Size of Hough Accumulator

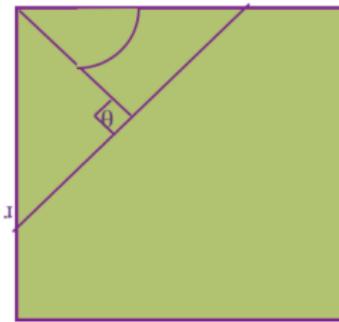
- ▶ The accumulator **A** becomes very large if:
 - ▶ Many parameters are used
 - ▶ Parameters are allowed to take many values
- ▶ The accumulator can be **much larger** than the image!
- ▶ It is practical to implement the accumulator **A** as a **single vector**, otherwise many matrix entries may always be **empty**, taking up valuable space

Accumulator Design

- ▶ The **design** of the Hough accumulator **A** is critical to keep its **dimensions** and **size** manageable
- ▶ Creating a manageable Hough accumulator is an art
- ▶ The following are **general** steps to follow

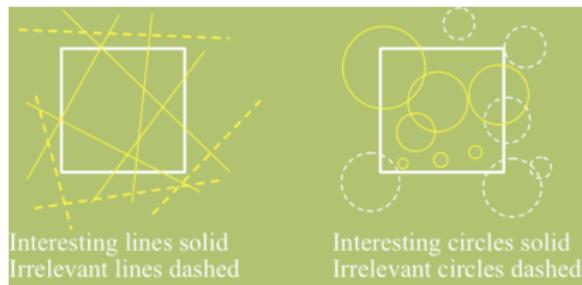
Accumulator Design – Step One

- ▶ Use **appropriate** curve equations. For example, in **line detection** the slope-intercept version is poor since **nearly vertical** lines have **large slopes**
- ▶ A better line representation: **polar form**
 $f(i, j; \mathbf{a}) = i\cos\theta + j\sin\theta - r = 0$, where $\mathbf{a} = (r, \theta) = (\text{distance, angle})$



Accumulator Design – Step Two

- ▶ **Bound** the parameter space to only allow the parameters for curves that **sufficiently intersect** the image
- ▶ What is “allowable” depends on the application – perhaps circles
 - ▶ Must lie **completely inside** the image
 - ▶ Must be of some **min** and **max** radius



Accumulator Design – Step Two Example

- ▶ In an $N \times N$ image indexed $0 \leq i, j \leq N - 1$, **detect lines** of the form $f(i, j; \mathbf{a}) = i\cos\theta + j\sin\theta - r = 0$ that **intersect** the image
- ▶ Thus each line must intersect the four sides of the image in **two places**
- ▶ **Either** the i -intercept or the j -intercept must fall in the range $0, \dots, N - 1$:
 - ▶ i -intercept $= r/\cos\theta$
 - ▶ j -intercept $= r/\sin\theta$
- ▶ So, can bound as follows: **either**
 $0 \leq r/\cos\theta \leq N - 1$ **or**
 $0 \leq r/\sin\theta \leq N - 1$

Accumulator Design – Step Three

- ▶ **Quantize** the accumulator space. Curves are **digital** so the accumulator is a **finite array**
- ▶ Decide how “finely-tuned” the detector is to be
- ▶ For example, for circles the choice may be easy – let circle centers (i_0, j_0) and circle radii be **integer values** only
- ▶ For lines, the selection can be more complex since digital lines aren’t always “straight”

Accumulator Design – Step Three Example

- ▶ In an $N \times N$ image, **detect circles** of the form $f(i, j; \mathbf{a}) = (i - i_0)^2 + (j - j_0)^2 - r^2 = 0$
- ▶ Radius and center constraints: $3 \leq r \leq 10$ and $r \leq i_0 \leq (N - 1) - r$ and $r \leq j_0 \leq (N - 1) - r$
- ▶ Assume radii and circle centers are **integers**: $r \in \{3, \dots, 10\}$ and for **each** r the circles are completely constrained within the image: $r \leq i_0 \leq (N - 1) - r$ and $r \leq j_0 \leq (N - 1) - r$
- ▶ Then the accumulator contains the following number of slots:
 $\# \text{ circles types} = [N - 2r]^2 = 8N^2 - 208N + 1520 \approx 8N^2$
- ▶ This is about 8 times the image size!
- ▶ Circle counting algorithms often assume only a few radii, example in coin counting only a few coin sizes are considered

Accumulator Design – Step Four

- ▶ Initialize accumulator **A** to 0
- ▶ At each **edge coordinate** (i, j) , $[E(i, j) = 1]$, increment **A**:
for all accumulator elements **a** such that (ϵ small)
 $|f(i, j; \mathbf{a})| \leq \epsilon$ set $\mathbf{A}(\mathbf{a}) = \mathbf{A}(\mathbf{a}) + 1$ until all edge points have
been examined
- ▶ **Threshold** the accumulator. Those parameters **a** such that
 $\mathbf{A}(\mathbf{a}) \geq \tau$ represent instances of the curve being detected
- ▶ Since local **regions** of the accumulator may fall above
threshold, detection of local **maxima** is usually done

Comments on the Hough Transform

- ▶ **Substantial memory** required even for simple problems such as line detection
- ▶ **Computationally intensive**; all accumulator points are examined and potentially incremented as each image pixel is examined
- ▶ **Refinements** of the Hough transform are able to deal with these problems to some degree

Refining the Hough Transform

- ▶ The best way to reduce **both** computation **and** memory requirement is to reduce the size of the Hough array. Here is a **very general** modified approach:
 1. **Coarsely quantize** the accumulator. Define **error threshold** ϵ_0 and **variable thresholds** τ_m, ϵ_m
 2. **Increment** the Hough array wherever ($\epsilon_m > \epsilon_0$)
 $|f(i, j; a)| \leq \epsilon_m$
 3. **Threshold** the Hough array using a threshold τ_m
 4. **Redefine** the Hough array by
 - ▶ Only allowing values similar to those $\geq \tau_m$
 - ▶ Requantize the Hough array more finely over those values
 5. Unless $\epsilon_m < \epsilon_0$, set $\epsilon_{m+1} < \epsilon_m, \tau_{m+1} < \tau_m$ and goto (2)
 6. The exact approach taken in such a strategy will vary with the application and available resources, but generally the Hough array can be made much smaller (hence faster)

Hough Transform – Unusual Circles

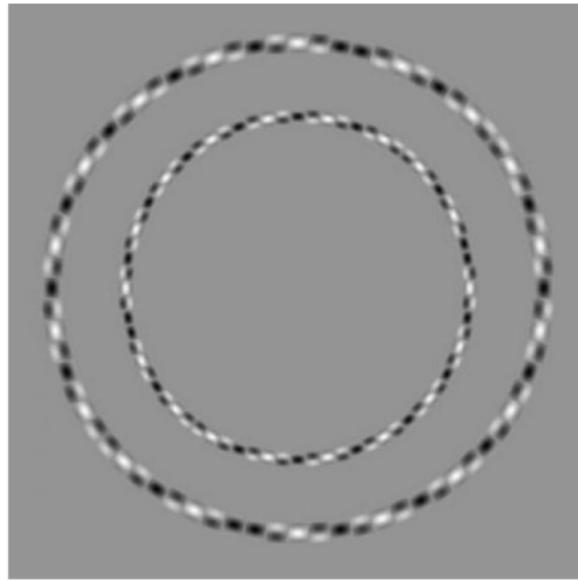


Figure: Could you find these using the Hough transform?

Template Matching

- ▶ Consider **finding instances** of a certain **sub-image**, example a certain object, written character etc
- ▶ This is called **visual search**
- ▶ Visual search is a **hard problem!**

Where's Waldo?



Waldo



Template Matching

- ▶ **Template matching** is a window-based technique to accomplish this task
- ▶ A **template** is a **sub-image**:

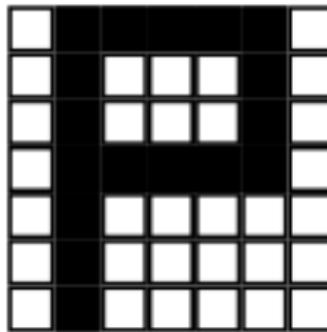


Figure: Template image of 'P'

Find The Faces

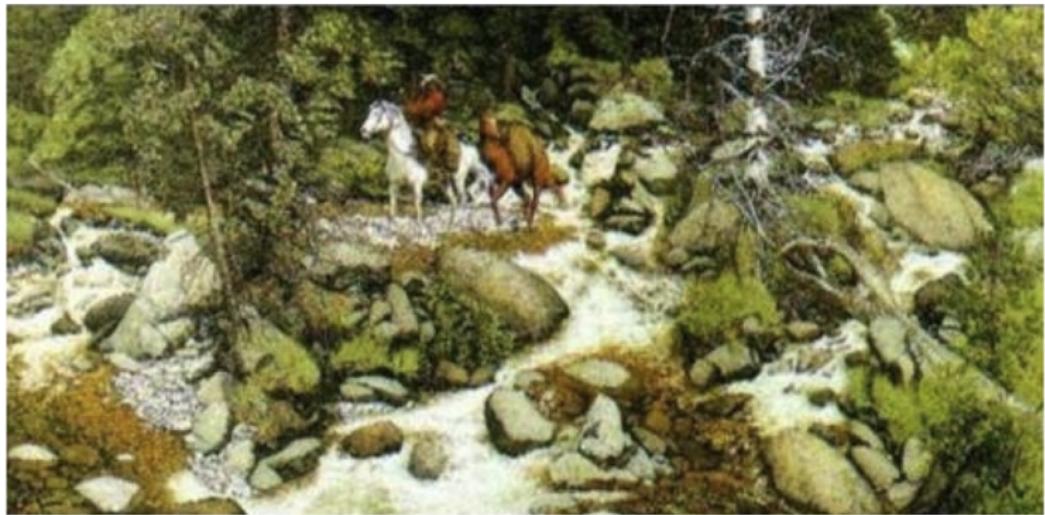


Figure: 11 faces!

Template

- ▶ Associated with a **template T** is a window \mathbf{B}_T which defines its domain:
$$\mathbf{T} = \{T(m, n); (m, n) \in \mathbf{B}_T\}$$
- ▶ We will also use window \mathbf{B}_T to search \mathbf{I}
- ▶ As before, the **windowed set** at (i, j) is:
$$\mathbf{B}_T \diamond \mathbf{I}(i, j) = \{I(i + m, j + n); (m, n) \in \mathbf{B}_T\}$$

Matching the Template to the Image

- ▶ Goal: Measure **goodness-of-match** of template \mathbf{T} with image patches $\mathbf{B}_{\mathbf{T}} \diamond \mathbf{I}(i, j)$ for all (i, j)
- ▶ Result: A **match image \mathbf{J}** taking large values where $\mathbf{B}_{\mathbf{T}} \diamond \mathbf{I}(i, j)$ and \mathbf{T} are highly similar
- ▶ We start with various **mismatch measures** and will derive a **match measure** from one of them

Mismatch Measures

- ▶ **Mismatch** $\{B_T \diamond I(i, j), T\}$
 1. Max absolute error: $\max_{(m,n) \in B_T} |I(i+m, j+n) - T(m, n)|$
 2. Mean absolute error: $\frac{1}{MN} \sum_{(m,n) \in B_T} |I(i+m, j+n) - T(m, n)|$
 3. Mean squared error: $\frac{1}{MN} \sum_{(m,n) \in B_T} |I(i+m, j+n) - T(m, n)|^2$
- ▶ All three are reasonable mismatch criteria
- ▶ Only (3) leads to an easy **match measure**

Mismatch Measures

- ▶ Decompose MSE into **three terms**:

$$\sum_{(m,n) \in \mathbf{B}_T} \sum [I(i+m, j+n)^2] -$$

$$2 \sum_{(m,n) \in \mathbf{B}_T} \sum I(i+m, j+n) T(m, n) + \sum_{(m,n) \in \mathbf{B}_T} \sum [T(m, n)]^2 \\ = E_{\mathbf{B}_T \diamond \mathbf{I}(i,j)} - 2C_{\mathbf{B}_T \diamond \mathbf{I}(i,j), \mathbf{T}} + E_T$$

- ▶ MSE is **small** when the match is good. Only the correlation term contributes to the match measure

Cross-Correlation Matching

► Schwarz Inequality

$$\sum_{(m,n)} A(m,n)B(m,n) \leq \sqrt{\sum_{(m,n)} [A(m,n)]^2 \sum_{(m,n)} [B(m,n)]^2}$$

- Equality holds iff $A(m,n) = KB(m,n)$ for all (m,n) where K is **any constant**

Upper Bound on Cross-Correlation Matching

- ▶ Upper bound the cross-correlation:

$$\begin{aligned} C_{\mathbf{B}_T \diamond I(i,j), T} &= \sum_{(m,n) \in \mathbf{B}_T} I(i+m, j+n) T(m, n) \\ &\leq \sqrt{\sum_{(m,n) \in \mathbf{B}_T} [I(i+m, j+n)]^2 \sum_{(m,n) \in \mathbf{B}_T} [T(m, n)]^2} \\ &= \sqrt{E_{\mathbf{B}_T \diamond I(i,j)} \cdot E_T} \end{aligned}$$

- ▶ **Equality** holds iff $I(i+m, j+n) = K T(m, n)$ for all $(m, n) \in \mathbf{B}_T$

Normalized Cross-Correlation

- ▶ Cross-correlation is a **good match measure** but we need to compare it with the theoretical upper bound: $\sqrt{E_{B_T \diamond I(i,j)} \cdot E_T}$ or normalize with respect to it
- ▶ Define $\hat{C}_{B_T \diamond I(i,j), T} = \frac{C_{B_T \diamond I(i,j), T}}{\sqrt{E_{B_T \diamond I(i,j)} \cdot E_T}}$
- ▶ $0 \leq \hat{C}_{B_T \diamond I(i,j), T} \leq 1$ for all (i, j)
- ▶ Define the **normalized cross-correlation image**:
 $J = \text{CORR}[I, T, B_T]$
- ▶ $J(i, j) = \hat{C}_{B_T \diamond I(i,j), T}$ if $0 \leq i \leq N - 1, 0 \leq j \leq M - 1$

Thresholding

- ▶ A good way to compute the **best** match is to find the **largest value** in the match image
 $J = \text{CORR}[I, T, B_T]$
- ▶ This assumes there is **exactly one match**. Alternately, **threshold** the match image:
 $K(i,j) = 1 \text{ if } J(i,j) \geq \tau$
- ▶ This will hopefully identify all points having a sufficiently good match
- ▶ If $\tau = 1$, the only **perfectly matches** will be found. Rarely happens due to noise
- ▶ τ usually chosen to be **close to** but **less than** one
- ▶ Found by trial and error

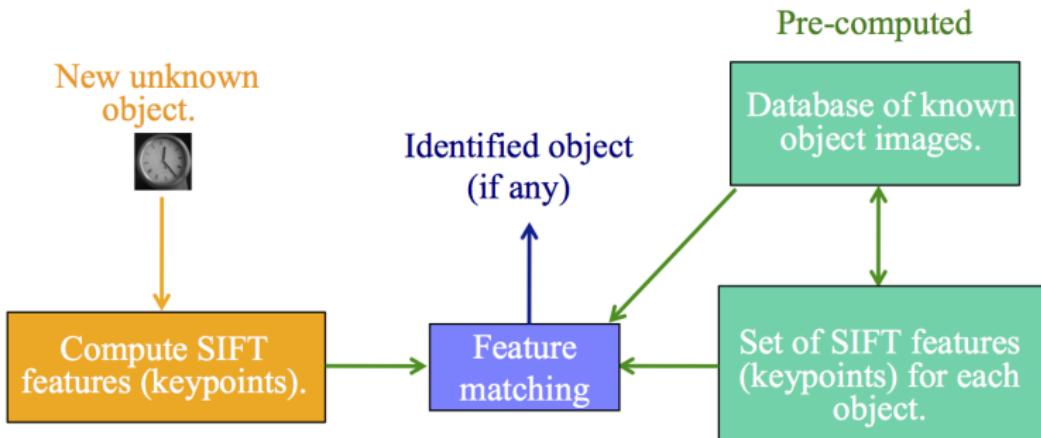
Limitations of Template Matching

- ▶ Template matching is quite **sensitive to object variations**
- ▶ Noise, rotation, scaling, stretching, occlusion and many other changes can confound the matching process

Scale Invariant Feature Transform

- ▶ **SIFT** (Scale Invariant Feature Transform) is a powerful and popular method of **visual search/object matching**
- ▶ Brainchild of Prof. David Lowe of University of British Columbia
- ▶ It can find objects that have been **rotated** and **scaled**
- ▶ **Limited ability** to match **objects** that have been **deformed**
- ▶ We will outline the **steps involved** and the **features** that are used in SIFT

SIFT – Basic Idea



The number of “keypoint” SIFT features may be large. They are related to edges.

SIFT Features

- ▶ Keypoint **candidates** are the maxima of **DoG-filtered responses**. Given gaussian at **scale** σ :

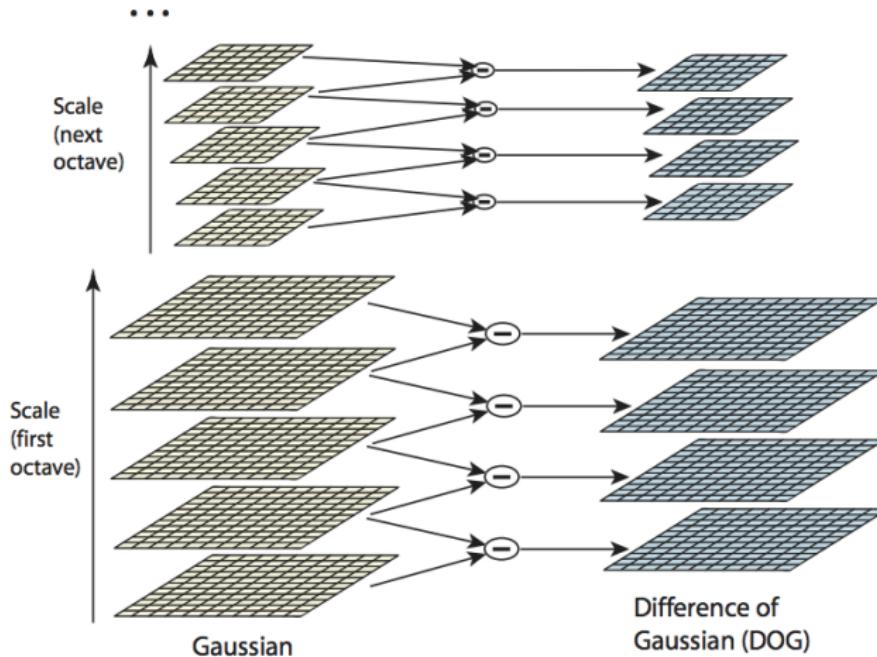
$$G_\sigma(i, j) = \frac{1}{2\pi\sigma^2} \cdot \exp\{-[i^2 + j^2]/2\sigma^2\}$$

- ▶ **DoG response** to image I is:

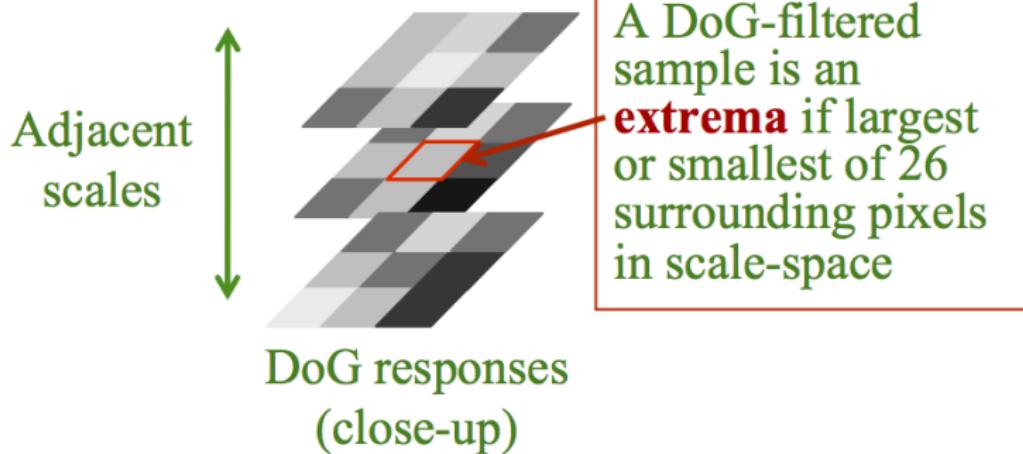
$$D_\sigma(i, j) = [G_{k\sigma}(i, j) - G_\sigma(i, j)] * I(i, j) \text{ where } k = 2^{(1/s)} \text{ and often } s = 2$$

- ▶ Note: $G_{k\sigma}(i, j) - G_\sigma(i, j) \approx (k - 1)\sigma\nabla^2 G_\sigma(i, j)$

DoGs Computer Over Scale



Extrema Detection



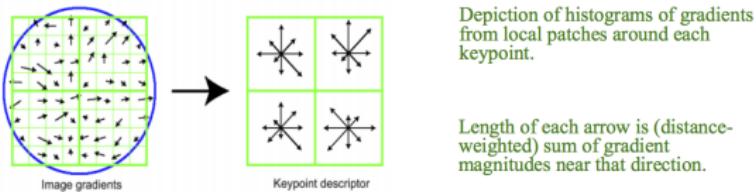
All **extrema** are found over **all DoG scales**

Extrema Processing

- ▶ Each “extrema” keypoint is then **carefully examined** to see whether:
 - ▶ It is actually **located on an edge**
 - ▶ If so, if the edge has a **strong enough contrast**
- Otherwise it is **rejected**
- ▶ Remaining extrema are **assigned an orientation** that is the gradient orientation of the **DoG response**

SIFT Features at Keypoints

- ▶ All **keypoints** for an image stored in the database and associated with **location, scale, and orientation**

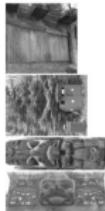


- ▶ At the **location** and **scale** of each keypoint, the **gradient magnitudes and orientations (relative to the keypoint orientation)** are found at all points in a patch around the keypoint
- ▶ These are formed into **descriptors** that vectorize the **histograms** of the gradient magnitudes/orientations from sub-patches of the larger patch
- ▶ Lastly, these vectors are made **unit-length** to reduce the effects of illumination

Matching

- ▶ When a new image is to be matched, the **SIFT descriptors** are computed from it
- ▶ The database consists of a set of images, possibly large, each with **associated SIFT descriptors**
- ▶ **Matching the image descriptors to the database** requires a process of **search** which can be done in many possible ways. Lowe uses **nearest-neighbor** search (Euclidean distance)
- ▶ Lowe also uses a **Hough Transform**-like method to **cluster** and **count keypoint descriptors**. This matching process also uses keypoint location, orientation, and scale **relative** to those found in the database
- ▶ Lowe's paper has more details

Lowe's Example



Recognition results

Outer boundaries
show matches and
keypoints are
squares

Comments on SIFT

- ▶ SIFT uses both **fundamental** and **ad-hoc** ideas. Yet, works very well within its domain (finding rigid bodies)
- ▶ It has **strong invariance** to:
 - ▶ Rotations
 - ▶ Translations
 - ▶ Scale changes
 - ▶ Illumination changes
- ▶ It has **reasonable invariance** to:
 - ▶ Occlusions
 - ▶ Affine transformations
- ▶ It has **weak invariance** to:
 - ▶ Objects that deform or change (like faces)